


Block size and GPU memory


NVIDIA L40 · XOR input · bitshuffle · fixed chunk and batch geometry · 6 GiB allocation budget

— LZ4 — Zstd ▲ observed median delta - - - explicit allocation estimate

256 KiB Zarr chunks · 144 MiB padded batch



1 MiB Zarr chunks · 288 MiB padded batch



Observed usage is a before/after free-memory delta while the stream is alive; it is not a sampled peak.

The baseline follows CUDA context creation. Runtime overhead is separate from explicit allocations.